

Software

Athanasios Kousathanas, Pablo Duchen & Daniel Wegmann

Department of Biology, University of Fribourg, Switzerland

March 15, 2015

1 Introduction

There are currently many programs available to conduct ABC analyses (see Table 1). However, most programs are specific to a particular problem, and the large majority for questions that typically arise in a population-genetics setting. These include programs to infer demographic histories (e.g. *ONeSAMP*, *PopABC*, *msABC* or *DIYABC*), to infer F-statistics (*ABC₄F*) or to infer parental contributions in an admixture event (*2BAD*). However, there exist also programs for phylogeographic inference (*msBayes*), Systems Biology (*ABC-SysBio*) or the inference under stochastic differential equations (the MATLAB package *abc-sde*).

A major benefit of such specific programs is that a model parameterization suitable for simulation-based inference has already been worked out and a set of summary statistics informative for the problem identified. However, a particularly powerful aspect of ABC is its application to almost any inference problem and model, and we will thus focus here on general purpose ABC software designed to be helpful in a large array of ABC applications, either to conduct the whole analysis pipeline or at least specific parts of it.

Specifically, we will discuss two similar ABC pipelines, one provided through the command line program *ABCtoolbox* (version 2.0, Wegmann et al. [2010]), and the other by means of combining the two R packages *abc* (version 2.0, Csilléry et al. [2012]) and *EasyABC* (version 1.4, Jabot et al. [2013]). Both pipelines offer very similar features and build around a similar logic: 1) they both provide utilities to generate a large set of simulations using external simulation software using a variety of sampling algorithms, 2) they both offer algorithms to infer parameters from such a set of simulations, 3) they provide tools to conduct model choice from such sets of simulations performed under different models, and finally, 4) they both offer a series of functions to validate estimations as well as model choice. However, these pipelines differ in specific implementations of some algorithms, which we will outline below, as well as the general way users interact with them. The packages *EasyABC* and *abc* are used within the statistical software environment R and are hence well suited for people writing their simulation programs in R or for those familiar with the handling of data sets in this environment. In contrast, *ABCtoolbox* is written purely in C++ and run from the command line, resulting in generally faster execution and making it particularly well suited when command-line programs are used to conduct simulations or calculate summary statistics.

Table 1: General and specific purpose ABC software.

Software	Purpose	Reference
<i>ABCtoolbox</i>	General	Wegmann et al. [2010]
<i>abc</i>	General	Csilléry et al. [2012]
<i>ABC-EP</i>	General	Barthelmé and Chopin [2011]
<i>ABC_distrib</i>	General	Beaumont et al. [2002]
<i>ABCreg</i>	General	Thornton [2009]
<i>EasyABC</i>	General	Jabot et al. [2013]
<i>DIYABC</i>	Population genetics	Cornuet et al. [2008]
<i>msABC</i>	Population genetics	Pavlidis et al. [2010]
<i>ONeSAMP</i>	Population genetics	Tallmon et al. [2008]
<i>PopABC</i>	Population genetics	Lopes et al. [2009]
<i>2BAD</i>	Population genetics	Bray et al. [2010]
<i>ABC4F</i>	Population genetics	Foll et al. [2008]
<i>msBayes</i>	Phylogeography	Hickerson et al.
<i>ABC-SysBio</i>	Systems Biology	Liepe et al. [2010]
<i>abc-sde</i>	Stochastic differential equations	Piccini (2014)

In the remaining of this chapter, we will walk the reader through the general usage of these two ABC pipelines. In order to give the reader a chance to replicate our analysis easily, we will use rather simple models and always give a detailed description of all settings used, along with all the code required to replicate the described analyses. We will begin with parameter inference, as we believe this is the step for which the programs discussed here will be used most often. However, we will also discuss how to use these pipelines to perform model choice, and how to conduct simulations using existing software.

2 Toy models

We will introduce the usage of the program *ABCtoolbox* and the R packages *EasyAbc* and *abc* through their application to the problem of inferring the mean (μ) and variance (σ^2) of a normal (Model A) and a uniform (Model B) distribution from a random sample. We will then further show how to use these ABC pipelines to distinguish between these models using model choice.

Using such simple models has two major benefits. First, it will allow us to compare the ABC estimates with those obtained from full likelihood solutions. Second, it is straightforward and quick to generate data under these models using just a few lines of code, for instance when using the free statistical programming language R, which we will do here. However, note that we will also discuss how to generate simulations with existing programs in Section 4.4.

Table 2: Observed statistics of test data set

mean	var	median	min	max	range	Q1	Q3
0.1020954	1.136706	0.07878515	-2.021652	3.158866	5.180517	-0.5975965	0.7986314

2.1 Observed data and summary statistics

We will begin by generating a data set of 100 random samples under Model A (the normal distribution) for which we would like to infer parameters. This is readily done in R as follows:

```
sampleSize <- 100;
data.obs <- rnorm(sampleSize, mean=0, sd=1);
```

Note that for this test example we do know the true parameters $\theta = (\mu, \sigma^2)$ that we want to estimate. This will allow us to test the accuracy of the estimation. Also, note that we saved the sample size in a variable that we will use later when generating simulations with the same sample size.

Next, we will define a function to calculate summary statistics both on the observed as well as simulated data sets. As summary statistics, we will use here the sample mean, variance and median, along with the smallest (min), largest (max) value, the range (max - min) and the first and third quartile (Q1 and Q3). A vector containing these summary statistics is readily generated in R using the following function:

```
calc.stats <- function (x){
  S <- c(mean(x), var(x), median(x), range(x), max(x)-min(x),
    quantile(x, probs=c(0.25, 0.75)));
  names(S) <- c("mean", "var", "median", "min", "max", "range", "Q1", "Q3")
  return(S);
}
```

This will now allow us to calculate these summary statistics on the observed data set (object data.obs), save them in the variable S_{obs} (object S.obs) and to also save them in a file called normal.obs.

```
S.obs <- calc.stats(data.obs);
write.table(t(S.obs), file="normal.obs", quote=F, row.names=F);
```

While each run will produce different values, we provide the values we obtained and will use in the following in Table 2 to allow for the replication of all our results.

2.2 Generating simulations

We will next generate a large number (10,000) of simulations with parameter values drawn from prior distributions and calculate the associated summary statistics for each simulation. In order to allow for a direct comparison between the models, we will assume uniform prior distributions for the mean $\mu \sim U[-1, 1]$ and variance $\sigma^2 \sim U[0, 4]$ for both models. Simulations for the normal model (Model A) are then generated as follows:

```

nsim <- 10000;
P.normal <- data.frame(mu=runif(nsim, min=-1, max=1), sigma2=runif(nsim, min=0, max=4));
S.normal <- data.frame(matrix(data=0, ncol=length(S.obs), nrow=nsim));
names(S.normal) <- names(S.obs);
for ( i in 1:nsim ) {
  y <- rnorm(sampleSize, mean = P.normal$mu[i], sd=sqrt(P.normal$sigma2[i]));
  S.normal[i,] <- calc.stats(y);
}
write.table(cbind(P.normal, S.normal), file="simNorm.txt", quote=F, row.names=F);

```

Again, we saved the simulations also in a text file (`simNorm.txt`) in order to use them with *ABCtoolbox*. Note that *ABCtoolbox* requires the parameters and statistics in the same file, which is achieved by binding the data frames containing the parameters (`P.normal`) and statistics (`S.normal`) together using `cbind()`.

Generating simulations under the uniform model (Model B) is achieved similarly. However, since the R function `runif()` requires the two limits of the uniform distribution, we need to calculate them from our parameters μ and σ^2 .

```

P.unif <- data.frame(mu=runif(nsim, min=-1, max=1), sigma2=runif(nsim, min=0, max=4));
S.unif <- data.frame(matrix(data=0, ncol=length(S.obs), nrow=nsim));

for ( i in 1: nsim ) {
  y <- runif(sampleSize, min=P.unif$mu-sqrt(3*P.unif$sigma2), max=P.unif$mu+sqrt(3*P.unif$sigma2));
  S.unif[i,] <- calc.stats(y);
}
write.table(cbind(P.unif, S.unif), file="simUnif.txt", quote=F, row.names=F);

```

3 Parameter inference

The estimation of posterior distributions is straightforward with both *ABCtoolbox* and the R package *abc*. As a common feature, they both implement the basic rejection algorithm originally introduced by Tavaré et al. [1997] and Pritchard et al. [1999], but they differ in the post-sampling adjustment algorithms they offer. Specifically, the package *abc* implements the original post-sampling adjustment based on a local linear regression initially introduced by Beaumont et al. [2002], as well as an extension to non-linear models with heteroscedastic variance [Blum and François, 2010]. In contrast, *ABCtoolbox* offers an implementation of the ABC-GLM algorithm introduced by Leuenberger and Wegmann [2010]. In this section, we will discuss differences between these algorithms as well as how to use them in the present example.

3.1 Rejection Algorithm

abc To start using *abc*, the package has first to be installed and loaded in R, which is done by the following two commands:

```
install.packages("abc");  
library(abc);
```

Note that while the installation has to be done only once on a system, the package has to be reloaded every time a new R session is started.

To now conduct an ABC rejection on the data simulated under the normal model (Model A), simply use the function `abc()` with the argument `method="rejection"` and by specifying the tolerance to be applied.

```
rejection <- abc(S.obs, P.normal, S.normal, tol=0.01, method="rejection");
```

Here, `S.obs`, `P.normal` and `S.normal` refer to the vector of observed summary statistics S_{obs} and the data frames containing the simulated parameters and summary statistics, respectively, as generated under Section 2. The additional argument `tol` specifies the fraction of simulations to be retained based on their distance to the observed summary statistics. A tolerance of 0.01, for example, indicates that the posterior density will be estimated from the parameter values of the 1% of all simulations that produced summary statistics closest to the observed summary statistics based on an euclidean distance metric.

The package *abc* offers an internal plotting function `hist.abc()` to display posterior distributions. Since this function overloads the basic `hist()` function of R, it can be called on an `abc` object by simply typing

```
hist(rejection);
```

Alternatively, it is also possible to use general R functions such as `hist()` or `density()` to plot posterior distributions. The results we obtained for the observed data shown above is plotted using `density` in Figure 1.

ABCtoolbox In contrast to the R packages discussed here, the program *ABCtoolbox* is a program to be used from the command line, preferentially in a Unix environment. While it thus run from the Windows command prompt, it is recommend to use the *cygwin* environment on a Windows computer to benefit from all features of *ABCtoolbox*.

ABCtoolbox accepts input settings both directly from the command line, as well as through an input file. A list of all arguments relevant for estimation and discussed in this chapter is provided in Table 3. To perform parameter estimation for the normal distribution example, for instance, the following input file may be used:

```
task estimate  
simName simNorm.txt  
obsName normal.obs  
params 1-2
```

```
maxReadSims 100000
tolerance 0.01
maxCor 1.0
```

Here, the argument `task` is set to `estimate` in order to run *ABCtoolbox* in estimation mode. Then, the argument `simName` specifies the name of the file containing the performed simulations. This file is requested to contain the used model parameter values together with the associated statistics, the names of which are provided in the first line. Similarly, the file specified with the argument `obsName` should contain the summary statistics S_{obs} calculated from the observed data, again with the first line of the file containing the names of the statistics and the second line the associated values. Using the argument `params`, *ABCtoolbox* is further told which columns of the simulation file contain the model parameters to be estimated. Note that the simulation file may contain an arbitrary number of additional columns that will be ignored if they are neither specified to be model parameters with the argument `params` nor summary statistics also present in the file with the observed summary statistics.

The additional required arguments `maxReadSims` and `tolerance` specify the maximum number of lines (simulations) that will be read from the simulation file, and the fraction of simulations to be retained in the rejection step, respectively. Finally, the argument `maxCor` specifies the maximal allowed correlation between summary statistics. We will discuss the issues with correlated statistics below, but have added this option here to avoid *ABCtoolbox* complaining about the presence of highly correlated statistics in our toy models.

To run *ABCtoolbox* with this input file (assuming it was saved under the name `estimate.input`), simply run in the command:

```
./ABCtoolbox estimate.input
```

Alternatively, the example can be run without using an input file by specifying the commands in the command line as:

```
./ABCtoolbox task=estimate simName=simNorm.txt obsName=normal.obs
params=1-2 maxReadSims=100000 tolerance=0.01 maxCor=1.0
```

The output of such a run is found in a series of files, the names of which begin with a prefix that can be set with the argument `outputPrefix`, a tag referring to its content, and a number referring to the data set and model for which the estimation has been conducted. While an exhaustive list of all output filename tags discussed in this chapter is given in Table 4, we will focus here on the file with tag `BestSimsParamStats`, which contains the retained simulations and is used to plot the rejection posterior.

The posterior estimates for μ and σ^2 obtained from the ABC-rejection algorithm are readily plotted in R using the `density()` function.

```
ABCrej <- read.delim("ABC_output_model0_BestSimsParamStats_Obs0.txt", sep="\t");
plot(density(ABCrej$mu));
plot(density(ABCrej$sigma2));
```

The results we obtained for the observed data shown above are plotted using `density` in Figure 1.

3.2 Post-sampling adjustments

Posterior distributions estimated with the rejection algorithm tend to be much broader than the true posterior distributions. This is shown for the normal model in Figure 1, but has been observed generally and is due to the often relatively large distance thresholds leading to parameter values resulting in summary statistics rather distant from S_{obs} to be accepted. Obviously, this loss of precision can be reduced by being more restrictive in accepting simulations, but this may require unrealistically computational efforts, particularly in more complex models.

An alternative is to correct for the effect of using large thresholds by exploiting the often simpler relationship between model parameters and summary statistics locally around the observed summary statistics. In a landmark paper, Beaumont et al. [2002] assume a linear relationship between model parameters and summary statistics locally among the retained simulations and proposed to use this relationship to project the parameter values of all retained parameter values to S_{obs} . More recently, Blum and François [2010] introduced an extension of this approach by fitting a non-linear, heteroscedastic model using neural networks. Both of these algorithms are implemented in the R package *abc*.

In contrast, *ABCtoolbox* offers an implementation of the ABC-GLM algorithm introduced by Leuenberger and Wegmann [2010] that estimates a local likelihood function instead of directly targeting the posterior distribution. While potentially slightly slower, this formulation is flexible in the choice of prior distributions and allows for model choice based on the marginal density. In practice, however, all mentioned post-sampling adjustment algorithms tend to give very similar results and the reader is advised to validate any estimation carefully in which these algorithms produce diverting estimates.

abc The two post-sampling adjustments implemented in the R package *abc* are used by simply choosing the appropriate `method` when calling the `abc()` function. There are three different methods available: `loclinear`, `ridge` and `neuralnet`, which correspond, respectively, to the classic regression adjustment introduced by Beaumont et al. [2002], a version of this algorithm using a ridge regression to deal with extensive collinearity among statistics, and the non-linear, heteroscedastic regression proposed by Blum and François [2010].

The following commands will perform posterior estimation using these algorithms on the toy model introduced above.

```
regression <- abc(S.obs, P.normal, S.normal, tol=0.1, method="loclinear");  
neural <- abc(S.obs, P.normal, S.normal, tol=0.1, method="neuralnet");
```

The built-in function `hist()` can then again be used to plot the estimated posterior distributions.

```
hist(regression);  
hist(neural);
```

Another function provided by the package *abc* is `plot.abc`, which can be used to plot the densities of the estimated posterior distributions together with additional, informative plots such as the prior distribution, the distribution of euclidean distances, and the residuals of the regression. Since this function overloads the standard R function `plot()`, it is simply used as follows:

```
plot(regression,param=P.normal);
plot(neural,param=P.normal);
```

Alternatively, the estimated posterior distributions can also be plotted using the R function `density`. For that purpose, one has to access specific elements of the object returned by the `abc()` function, namely the projected model parameter values as `adj.values` as well as their weights `weights`. The following R commands, for instance, plot the posterior densities obtained via the regression and neural-network adjustment for μ :

```
plot(density(regression$adj.values[,1], weights=regression$weights));
plot(density(neural$adj.values[,1], weights=neural$weights));
```

Posterior distributions plotted using these functions are compared to those obtained through other methods in Figure 1. Note that the object returned also contains the retained model parameter values in the element `unadj.values` that can be used to plot the rejection posterior distribution.

```
plot(density(regression$unadj.values))
```

ABCtoolbox When running *ABCtoolbox* in `estimation` mode, the ABC-GLM adjustment introduced by Leuenberger and Wegmann [2010] is performed automatically and the results available in the output file with tag `MarginalPosteriorDensities`. To plot the posterior estimates in R, simply load that file and use the function `density`.

```
ABCGlm <- read.delim("ABC_GLM_model1_MarginalPosteriorDensities_Obs1.txt");
plot(ABCGlm$mu, ABCglm$density);
plot(ABCGlm$sigma2, ABCglm$density.1);
```

3.3 Multidimensional posteriors

abc Apart from marginal posterior distributions, both the R package *abc* as well as *ABCtoolbox* are capable of estimating multidimensional posterior distributions. In the case of *abc*, however, the posterior densities have to be estimated using standard functions of R such as `kde2d()` for two-dimensional posterior distributions. In our toy model, for instance, using

```
posterior2d <- kde2d(regression$adj.values[,1], regression$adj.values[,2], n=100);
contour(posterior2d);
```

where, `regression$adj.values[,1]` represents the projected model parameter values for μ , `regression$adj.values[,2]` represents the projected model parameter values for σ^2 , and `n=100` specifies the number of marginal grid points to be used for the density estimation. These R commands will produce a plot similar to the one in figure 1C.

ABCtoolbox To generate multidimensional posterior densities on a grid, simply add the argument `jointPosteriors`, followed by the names of the model parameters for which the multidimensional posterior distribution is to be estimated. In addition, the number of marginal grid points has to be

Table 3: *ABCtoolbox settings for estimation.*

Setting type	Setting	Description
Basic	task	Task to be performed. Possible options are: simulate, estimate, findStatsModelChoice
	estimationType	Standard estimation performs the GLM approach Leuenberger and Wegmann [2010]. Possible options:
	params	Specify the parameter columns in the file that contains the simulations.
	simName	File containing simulations.
	obsName	File containing observed summary statistics.
	tolerance	Fraction of simulations to retain.
	numRetained	No. of simulations to retain.
	maxReadSims	Maximum number of read simulations.
	pruneCorrelatedStats	Remove statistics that are correlated, possible options 0 (retain) or 1 (remove).
	maxCor	Maximum acceptable correlation coefficient between statistics.
	outputPrefix	Prefix for output files.
	writeRetained	Indicate whether to write retained simulations which can be used to obtain ABC rejection posteriors.
Validation	standardizeStats	Standardize statistics.
	obsPValue	The number of retained data sets for testing how well the inferred GLM model fits the observed data in multi-dimensional space (section 3.4.1).
	tukeyPValue	The number of retained data sets for performing the Tukey test (section 3.4.1).
	modelChoiceValidation	The number of cross-validation replicates for validating model choice (section 4.2).
	randomValidation	The number of cross-validation replicates for random parameter validation (section 3.4.3).
Posterior density	retainedValidation	The number of cross-validation replicates for retained parameter validation (section 3.4.3).
	posteriorDensityPoints	Number of points to estimate posterior density.
	diracPeakWidth	Smoothing parameter for posterior densities.
	jointPosteriors	Comma separated list of parameters for which the joint posterior is to be evaluated (section 3.3).
	jointPosteriorDensityPoints	No. of points to evaluate joint posterior (section 3.3).

Table 4: *ABCtoolbox* estimation output files.

File type	File tag	Content
Basic	MarginalPosteriorCharacteristics	Characteristics of marginal posterior distributions (e.g., mode, mean, quantiles).
	BestSimsParamStats	Retained simulations from ABC-rejection.
	MarginalPosteriorDensities	GLM-adjusted marginal posterior densities.
	jointPosterior	GLM-adjusted joint posterior estimates.
	modelFit	Model choice results including bayes factors and posterior support for compared models.
Parameter validation	RandomValidation	Results from random validation.
	RetainedValidation	Results from retained validation.
Model choice validation	modelChoiceValidation	Results for model choice validation.

specified using `jointPosteriorDensityPoints`. The joint posterior estimates for the parameters μ and σ^2 of the normal distribution model, for instance, is thus estimated by simply running *ABCtoolbox* with a modified `estimate.input` input file that contains these two additional lines:

```
jointPosteriors mu,sigma2
jointPosteriorDensityPoints 100
```

Note that this algorithm is not suitable to estimate densities in many dimensions as the total number of grid points grows exponentially with the dimensionality. In this above example, the density will be evaluated at $100 \times 100 = 10^4$ positions. Running such a command for a four dimensional posterior will already result in 10^8 positions at which the density has to be estimated. To generate samples from posterior distributions in high dimensions, *ABCtoolbox* also implements an MCMC algorithm to sample from the full joint posterior distribution. While we will not discuss that algorithm here, we refer the user to the manual of *ABCtoolbox* for more details.

When running *ABCtoolbox* with the additional arguments `jointPosteriors` and `jointPosteriorDensityPoints` the posterior density at each grid point will be written to an output file with tag `jointPosterior`. The resulting joint posterior of μ and σ^2 can then be plotted in R using the function `contour()`.

```
plot2D <- read.delim("ABC_GLM_model10_jointPosterior_1_2_Obs0.txt");
x <- unique(plot2D$mu);
y <- unique(plot2D$sigma2);
z_density <- matrix(data=plot2D$density,nrow=length(x),ncol=length(y),byrow=F);
contour(x,y,z_density,xlab=expression(mu),ylab=expression(sigma^2));
```

Since densities are often hard to interpret, *ABCtoolbox* also calculates and prints the smallest high posterior density interval (HDI) including each grid point to the same output file. The HDI corresponds to a posterior credible interval in the multidimensional parameter space and hence allows the generation of contour plots where contour lines indicate posterior credible intervals as follows.

```
z_HPD <- matrix(data=plot2D$HDI,nrow=length(x),ncol=length(y),byrow=F);
contour(x,y,z_HPD,xlab=expression(mu),ylab=expression(sigma^2));
```




Figure 1: Posterior densities for the mean (A) and variance (B) parameters of the normal distribution model produced by a variety of methods. Dashed grey lines indicate prior densities. (C) Joint posterior density for μ and σ^2 produced with *ABCtoolbox*. Dashed black lines indicate true parameter values for all panels.

The two dimensional posterior distribution we obtained this way for the normal distribution example is given in Figure 1C.

3.4 Validation of parameter estimation

3.4.1 Using a wrong model

An essential first validation step is to check whether the observed statistics can be reproduced by the examined model. A failure of the model to reproduce some of the statistics may indicate that a model is either not reflecting reality close enough, or that inappropriate prior distributions have been used (e.g., too narrow distributions). More importantly, all post-sampling adjustments assume that the model fitted to the model parameters and summary statistics can be used to either accurately project retained simulations to S_{obs} (the methods implemented in *abc*), or is an accurate description of the likelihood of S_{obs} with the parameter range of the retained simulations (the method implemented in *ABCtoolbox*). A violation of these assumptions leads to an extrapolation to an area of the summary statistics space for which no samples have been obtained, and is hence prone to biased inference.

Checking if the observed summary statistics S_{obs} are within the range of summary statistics generated by the model is, however, a bit tricky in higher dimensions. For instance, consider the marginal summary statistics distributions shown in Figure 2 for the normal and uniform toy models, respectively. These distributions were plotted in R using the `density()` function directly from the simulated data.

```
plot(density(S.normal$var), col='blue',);
lines(density(S.unif$var), col='red');
lines(rep(S.obs[2],2), par("usr")[3:4]);
```

These plots suggest that both summary statistics are readily generated by both models. However, a mismatch might manifest when looking at combinations of summary statistics. For this one may plot two-dimensional distributions of pairs of summary statistics using the R functions `kde2d` and

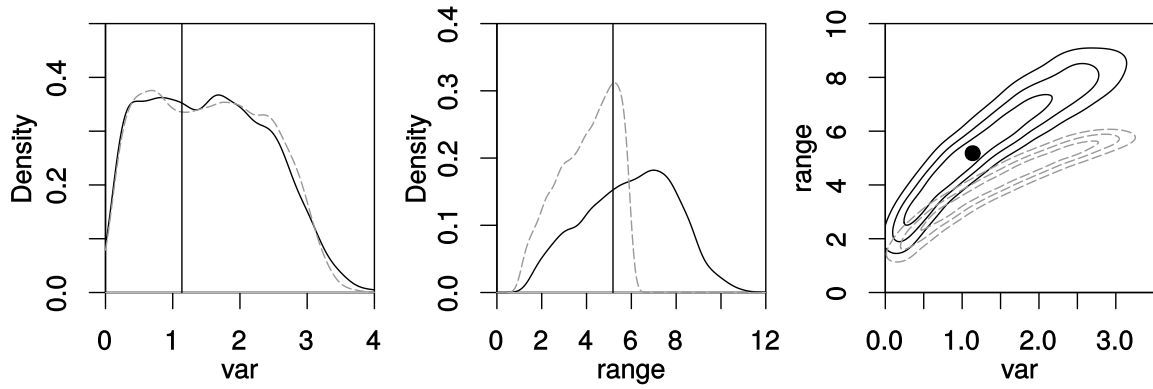


Figure 2: The distributions of simulated versus observed statistics *variance* (left panel) and *range* (middle panel) and their joint distribution (right panel) for the normal (blue) and uniform (red) distribution models. Observed data is shown with a black vertical line in the left and middle panels and with a black dot in the right panel.

contour.

```
d <- kde2d(S.unif$var, S.unif$range, n=100);
contour(d$x, d$y, d$z);
```

As is shown in the rightmost panel in Figure 2, the combination of the two observed summary statistics *variance* and *range* can indeed not be reproduced by the uniform model. However, this fact is not visible when looking at marginal densities only.

ABCtoolbox Since visual inspection is only fruitful for a limited number of dimensions, *ABCtoolbox* offers two statistical tests for assessing whether a given model can reproduce the observed data S_{obs} in the multidimensional statistics space. The first test compares the marginal density of the observed data to the marginal density of the retained simulations. The fraction of retained simulations with smaller or equal marginal density than the observed data is then provided as *the marginal density P-value* where small values indicated a poor fit of the model to the observed data.

The second test evaluates how central the observed data lies within the multidimensional cloud of retained simulations by comparing the Tukey depth (Cuesta-Albertos and Nieto-Reyes [2008]) of the observed data to the same depth of the retained simulations. The fraction of retained simulations with smaller or equal depth than the observed data is then reported as the *Tukey P-value*. Again, a low Tukey P-value indicates issues with generating simulations close to the observed data and hence suggests a poor fit of the model. However, note that the opposite is not necessarily true. Indeed, even a poor model that is capable of generating a cloud of summary statistics surrounding S_{obs} will pass both tests.

To perform these tests, simply call *ABCtoolbox* with the arguments `marDensPValue` and `tukeyPValue`, where each of them indicates the number of retained simulations to be used when calculating the respective P-value. When adding the following two lines to the input file `estimate.input`, for instance, *ABCtoolbox* will use 1,000 retained simulations to evaluate the P-values.

Table 5: Observed *P*-value and Tukey *P*-value results for normal distribution example

Model	marginal density	marginal density P-Value	Tukey Depth	Tukey P-Value
1	0.162	0.96	0.13	0.963
2	8.81×10^{-25}	0	0	0

```
marDensPValue 1000
tukeyPValue 1000
```

The results we obtained for these tests for our toy models is shown in Table 5. As expected from the visual inspection in Figure 2, the uniform model is not capable of reproducing the observed data and hence fails both tests.

3.4.2 Cross-validation / Accuracy of point estimates

The accuracy of posterior point estimates is generally done by estimating the parameters for data sets for which the true parameter values are known. This is readily done in an ABC setting as a leave-one-out test in which one of the provided simulations is randomly chosen and all other are used to infer the parameter estimates for this data (often called “pseudo-observed” data). The inferred posterior point estimates such as the maximum a posteriori (MAP or posterior mode), the posterior mean or the posterior median are then plotted against the parameter values used to generate the data (referred to as the “true parameters”). This process (also called cross-validation) is then repeated for many “pseudo-observed” data sets to obtain a general measure of accuracy. The procedure may also be repeated to test specific ABC settings such as the effect of the choice of tolerance or the number of available simulations.

abc To use this cross-validation algorithm with the R package *abc*, simply call the function `cv4abc()` with the arguments matching those of the estimation plus the additional argument `nvals`, which specifies how many pseudo-observed data sets are to be used. However, note that the observed data does not have to be provided, as it is not used in cross validation. The following code, for instance, will conduct cross validation on the normal distribution example for the neural network estimation algorithm based on 100 individual pseudo-observed data sets.

```
cv.neural <- cv4abc(P.normal, S.normal, tols=0.1, statistic= mode, method="
neuralnet", nval=100);
```

Such a call will return an object containing both the true parameter values (element `true`) as well as the estimated parameter values (element `estim`), which can be used to estimate correlations among them and plot for visual inspection. A plot such as the one shown in Figure 3 is generated by

```
plot(cv.neural);
```

ABCtoolbox *ABCtoolbox* offers two flavors of this cross-validation algorithm: either by picking simulations randomly, or by picking simulations only among those that were retained. The former, which is invoked through the argument `randomValidation`, corresponds to picking parameter values from the prior distribution and is thus informative above the overall accuracy of the ABC estimation under the chosen model. The latter, which is invoked using the argument `retainedValidation`, is informative about the accuracy of the estimation for the parameter space leading to similar data as the one observed.

When running *ABCtoolbox* with one or both of those arguments, an additional output file with tag `randomValidation` or `retainedValidation` is generated that contains the true parameter values along with a series of posterior point estimates for each parameter, namely the MAP (or mode) as well as the posterior mean and median. This file can then be loaded into R to generate plots such as those shown in Figure 3 as follows:

```
validation <- read.delim("ABC_GLM_model0_randomValidation_Obs0.txt");
plot(validation$mu, validation$mu_mode);
```

3.4.3 Checking for biased posteriors

Pseudo-observed data sets can be used equally to detect potential biases in the marginal posterior distributions. If the posterior distributions of a parameter were unbiased, the position of the true parameters across many replicates must be given by the posterior densities. We recently proposed to test this directly using the probability integral transform test (PIT histogram) Wegmann et al. [2009]. This is done by recording the position of the true parameter value in the cumulative posterior distribution (the posterior quantile) for each pseudo-observed data set. In case posteriors were unbiased, these posterior quantiles must be uniformly distributed between 0 and 1. Similarly, the smallest high posterior density intervals (HDI) containing the true parameter value must be distributed uniformly.

ABCtoolbox The procedure for this is the same as the one described in section 3.4.2. The same output file will contain information that allows us to check for biased posteriors. For instance, the output from these analyses can be used to determine whether the posterior quantiles and HDI are uniformly distributed either by visual inspection or by performing a statistical test such as the Kolmogorov-Smirnov test. The result of the random and retained validation analyses for the normal distribution example can be seen in Figure 4A and Figure 4B, respectively.

4 Model choice

While model choice is commonly used in Bayesian statistics, it is contentious when using ABC due to the problem that even sufficient summary statistics may lead to biased inferences [Robert et al., 2011]. Nonetheless, ABC model choice has been used successfully in practice and both ABC pipelines discussed here offer algorithms to conduct model choice when simulations from multiple models are

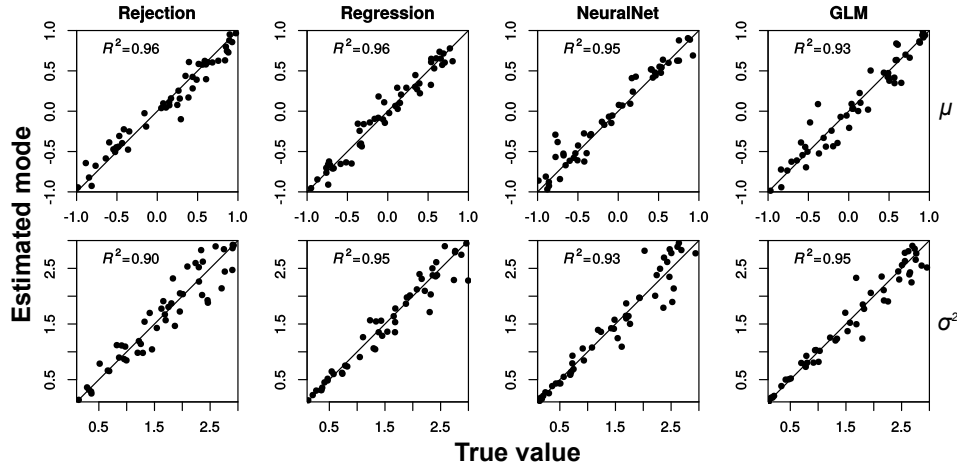


Figure 3: Parameter validation using different methods. The estimated posterior mode for mean and variance is plotted against the true values.

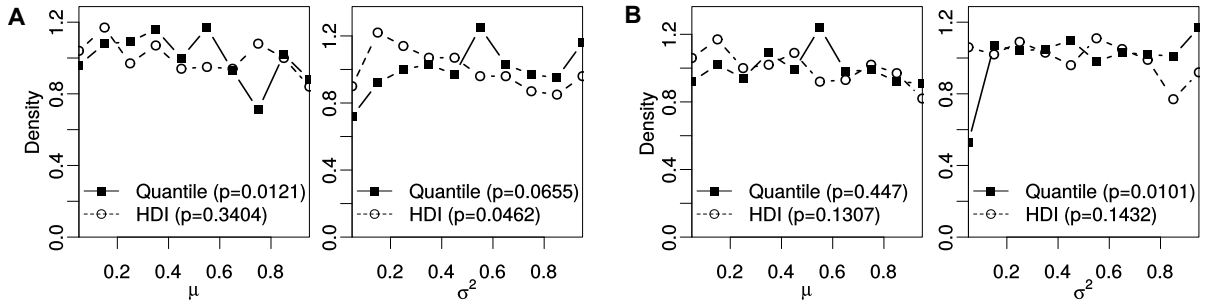


Figure 4: Parameter validation testing uniformity of the distribution of posterior quantiles and HDI when using random (A) and retained (B) simulations.

available. However, the user is advised to validate any ABC model choice carefully and we will discuss here tools provided by the R package *abc* as well as *ABCtoolbox* to aid in that crucial step.

4.1 Inferring Bayes Factors

Bayesian model choice relies on the estimation of Bayes factors or model posterior probabilities. In standard Bayesian statistics, these are estimated from the marginal densities of the compared models, where the marginal density of model i is defined as the integral of the likelihood function, weighted by the prior distribution:

$$\mathbb{P}(M_i|D) = \int \mathbb{P}(D|\theta_i, M_i)\mathbb{P}(\theta_i|M_i)d\theta_i$$

In case the likelihood function is not available for analytical evaluation, both posterior probabilities and Bayes factors can be estimated using ABC. The algorithm implemented in *ABCtoolbox*, for instance, fits a likelihood model to the retained data, and then uses this model to analytically calculate the marginal density for each model. In contrast, the algorithms available in the R package *abc* use the fact that the marginal density is proportional to the fraction of simulations that resulted in simulations close to the observed data S_{obs} when the simulations are generated according to the

Table 6: Results of model choice on toy models

Model	Posterior Probability		Bayes Factor	
	Normal	Uniform	Normal	Uniform
<i>abc</i> (rejection)	0.67	0.33	2.06	0.49
<i>abc</i> (neuralnet)	1.00	0.00	1.56×10^2	0.0
<i>ABCtoolbox</i>	1.00	0.00	1.84×10^{23}	5.43×10^{-24}

prior distribution. The posterior probabilities of the different models are thus estimated from the relative number of simulations being close to S_{obs} either from direct counting or through a regression adjustment similar to parameter inference. Note that it is crucial for both algorithms that the exact same summary statistics have been calculated under both models.

To illustrate the model choice algorithms implemented, we will attempt to estimate which of the two toy models introduced above (the normal and uniform model) was used to generate the observed data. We refer the reader to section 2 for more details on those models.

abc In order to conduct model choice with *abc*, the summary statistics of all models have to be concatenated into a single data frame or matrix. In addition, a vector indicating for each simulation the model under which it has been generated has to be created. For our toy models, this is simply achieved as follows:

```
allSimulations <- rbind(modelA[,3:10],modelB[,3:10]);
index <- rep("norm",dim(S.normal)[1], rep("unif",dim(S.unif)[1]));
```

The actual model choice is then conducted using the function `postpr`, which takes as arguments the observed summary statistics S_{obs} , the object containing the summary statistics for all models, and the index vector. In addition, the tolerance for the rejection step as well as the method for estimating Bayes factors has to be provided. In total, *abc* offers three such methods. The simplest is the the method `rejection`, which estimates posterior probabilities of the different models directly from the relative proportions of accepted simulations. The two other methods, `mnlogistic` and `neuralnet` attempt to correct for the often large tolerance values by estimating the relative densities of retained simulations at S_{obs} using either a multinomial logistic regression Beaumont [2008], Fagundes et al. [2007] or neural networks François and Guillaume [2011], respectively.

The following command will run model choice using the `neuralnet` method on our toy models and, using the function `summary()`, print the results to screen in a nice format.

```
model.choice <- postpr(s_obs, index=index, sumstat=allSimulations, tol=0.1, method=
  "neuralnet");
summary.postpr(model.choice);
```

The results for our toy models is shown in Table 6. As is expected from the observation that the uniform model fails to reproduce the observed summary statistics, the preferred model for this data is the normal model. However, note that the results for the rejection method, which is also run

by default when performing a `mnlogistic` or `neuralnet` estimation, is much less clear due to the relatively large tolerance applied here.

ABCtoolbox To perform model choice with *ABCtoolbox*, simply provide the arguments `simName` and `params` for multiple models using semicolons. To run model choice on our toy models, for instance, the input file `estimate.input` provided above is modified as follows:

```
simName simNorm.txt ;simUnif.txt
params 1-2;1-2
```

When running *ABCtoolbox* with such an input file, an additional file with tag `modelFit` is generated. This file contains the marginal densities, Bayes factors and posterior probabilities for each model. The results from this file obtained for our toy models is shown in Table 6, clearly indicating that the normal model is a much better fit.

4.2 Model choice validation

As was shown recently, model choice conducted with ABC may lead to biased or even wrong posterior probabilities, even if the summary statistics are sufficient for all models compared (see Robert et al. [2011]). Therefore, careful validation is a key and compulsory step of any ABC model choice analysis. An initial first test may be to evaluate the power of choosing the correct model by means of pseudo-observed data sets. Such a cross-validation, which is offered by both *abc* as well as *ABCtoolbox*, simply picks random simulations among those provided from both models, conducts model choice, and records how frequently the correct model was preferred. In addition, *ABCtoolbox* provides means to test for biases in the obtained posterior probabilities by comparing the ABC posterior probability (termed p_{ABC}) against those empirically expected ($p_{empirical}$) [Peter et al., 2010].

abc The R package *abc* contains the function `cv4postpr` to conduct cross-validation for model choice. This function randomly picks one simulation from the file containing all simulations, performs model choice using the chosen simulation as pseudo-observed data, and records which model obtained the highest posterior probability. This is then repeated many times to determine the confusion matrix. As an example, consider the following call to `cv4postpr` to conduct 100 such replicates on our toy models using the index vector created above. To then print the confusion matrix, one may use the function `summary()` and to obtain a graphical representation the built in `plot()` function as follows.

```
cv.model.choice <- cv4postpr(index,allSimulations , nval=100 , tols=0.1 , method=
  neuralnet)
summary(cv.model.choice);
plot(cv.model.choice);
```

ABCtoolbox To perform model choice validation with *ABCtoolbox*, simply add the argument `modelChoiceValidation` followed by the number of pseudo-observed data sets to be used to the estimation file. To use 1000 pseudo-observed data sets, for instance, you may add the following line to the input file:

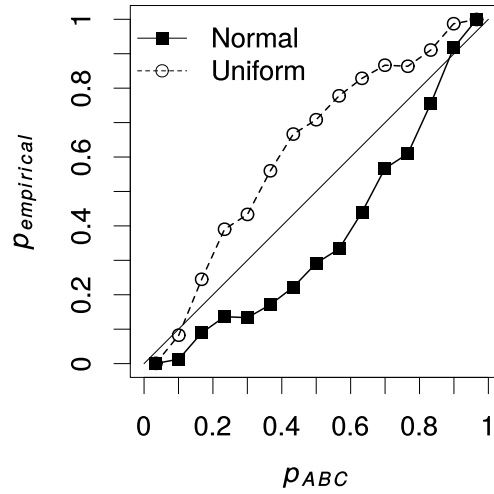


Figure 5: Posterior probability for normal and uniform distribution models estimated by *ABCtoolbox* (p_{ABC}) versus an empirical estimate of the same probability through simulation ($p_{empirical}$).

```
modelChoiceValidation 1000
```

ABCtoolbox will then perform cross validation and write the results to two different files. The first has the tag `confusionMatrix` containing the confusion matrix as well as statistics calculated from it. For the toy model, for instance, we learn from this file that the normal model is correctly identified from data generated under that model in $> 99\%$ of the cases.

The second file with tag `modelChoiceValidation` contains the raw results from the model choice validation and can be used for a more detailed validation analyses. For instance, we have recently proposed to compare the estimated model posterior probabilities with the empirical ones, an analysis that can reveal biases in ABC model choice [Peter et al., 2010, Chu et al., 2013]. The basic logic of this analysis is that among all pseudo-observed data sets that resulted in an ABC posterior probability $p_{ABC} = x$ in favor of, say, model 1, a fraction x should have been generated under model 1. To test for this, data sets are binned according to their ABC posterior probabilities and the fraction among them that were generated with model 1 is determined.

The *ABCtoolbox* package includes the Rscript `Make_Model_choice_power_plot.r` to conduct this analysis and to produce the plot shown in Figure 5. For our toy models it appears that there is a slight bias towards the normal model. This is evident from the fact that when $p_{ABC} = 0.5$, the data sets were actually generated from the uniform model in about 70% of the cases. However, among the data sets that resulted in very high p_{ABC} (≥ 0.99), the vast majority was generated under that model. Therefore, we have high confidence in the model choice results of our observed data, which produced ≥ 0.99 posterior probability support for the normal distribution model (Table 6).

4.3 Choosing summary statistics

4.3.1 Statistics for parameter inference

The choice of summary statistics is crucial in any ABC inference in that too few summary statistics are likely to miss out on important information and too many introducing harmful noise to the estimation [Wegmann et al., 2009, Beaumont, 2008, Blum et al., 2013]. To date, many methods have been proposed to choose informative summary statistics from a larger set (see Blum et al. [2013] for a review), but we will focus here on those available through the *ABCtoolbox* package, in particular the use of linear combinations of summary statistics.

The use of such linear combinations was first introduced by Wegmann et al. [2009], who proposed to find them by means of Partial Least Squares (PLS) regression. Broadly speaking, PLS is similar in spirit to a Principal Component Analysis (PCA), but instead of finding linear combinations that maximize the variance explained in the summary statistics space, PLS components are chosen such that they maximize the product of the variance among summary statistics and the covariance between parameters and statistics [Tenenhaus et al., 1995]. Recently, alternative means of finding linear combinations of summary statistics have been proposed, such as through boosting [Aeschbacher et al., 2012a] or by regressing summary statistics on to posterior means inferred from an initial set of simulations [Fearnhead and Prangle, 2012].

While all these methods are readily used with *ABCtoolbox* once the linear combinations have been found, we will illustrate the usage of this functionality based on the PLS approach, which is easy to implement with the R-package 'pls'. In fact, the *ABCtoolbox* package provides an R script to perform this analysis taking as input the simulations file (`simNorm.txt`). Performing a PLS analysis on the simulations from the normal distribution example reveals that 2 PLS components are sufficient for explaining the variance of the parameters of the normal distribution (Figure 6). This result is expected since the mean of a sample and the mean and variance of a sample are sufficient statistics for estimating, respectively, the mean and variance parameters of a normal distribution.

Any definition of linear combinations resulting from such a PLS or any other approach can then be used to transform the statistics of a set of simulations and the observed data using *ABCtoolbox*, and then used in parameter inference. The PLS R script mentioned above, for instance, writes the resulting PLS components to the file `PLSdef.txt`, which is then provided to *ABCtoolbox* and run in the `transform` mode as follows:

```
./ABCtoolbox task=transform linearComb=PLSdef.txt input=simNorm.txt
               output=simNorm.pls numLinearComb=2
./ABCtoolbox task=transform linearComb=PLSdef.txt input=normal.obs output=normal.obs.
               pls numLinearComb=2
```

Note that while we provided all arguments to *ABCtoolbox* on the command line, they may equally well be given in an input file.

Using the transformed summary statistics in the estimation step is then straightforward: simply provide the transformed files using the arguments `simName` and `obsName`.

As mentioned above, using alternative ways to find linear combinations is compatible with *ABC-*




Figure 6: Number of PLS components versus root mean square error of prediction (RMSEP) for the normal distribution model.

toolbox, as long as the linear combinations can be written in a definition file as the one created by the PLS script. Alternatively, the statistics may also be transformed with different software and then provided to *ABCtoolbox* (or the R package *abc*) for the estimation step.

One issue with using linear combinations of summary statistics is that information that arises from non-linear combinations of statistics are not taken into account. In such situations, it may be beneficial to increase the summary statistics space through combinations of summary statistics before finding linear combinations [Aeschbacher et al., 2012b]. As outlined below, *ABCtoolbox* has an option to also generate all pairwise products of summary statistics when generating simulations for this purpose.

4.3.2 Statistics for model choice

Finding appropriate combinations of statistics for model choice is particularly tricky. Just as for parameter inference, too few statistics may fail to capture important information, while too many are likely adding non-informative noise leading to biases in the inference. In addition, and as recently shown by Robert et al. [2011], even summary statistics sufficient for parameter inference may lead to biased Bayes factors. Consider two models \mathcal{M}_1 and \mathcal{M}_2 of shared parameters θ . If a set of summary statistics S was sufficient for both models, then the likelihood of the summary statistics and the likelihood of the full data are proportional for both models

$$\mathbb{P}(D|\theta, \mathcal{M}_1) = c_1 \mathbb{P}(S|\theta, \mathcal{M}_1)$$

$$\mathbb{P}(D|\theta, \mathcal{M}_2) = c_2 \mathbb{P}(S|\theta, \mathcal{M}_2)$$

However, there is no guarantee that the two proportionality constants c_1 and c_2 are identical, which leads to the Bayes factors that are off by c_1/c_2 .

Table 7: Combinations of statistics sorted by estimated power to distinguish models.

Rank	Power	Largest Pairwise Correlation	No.statistics	Statistics
1	1	0.773646	4	S1,S2,S4,S5
11	1	0.970552	3	S2,S6,S8
42	0.997	0.970552	2	S2,S6
111	0.732	0	1	S6

Unfortunately, the methods introduced above for finding good summary statistics for parameter inference are not readily extended to the problem of model choice. If one aims for using linear combinations, the most obvious choice is Linear Discriminant Analysis (LDA), as was recently proposed by Estoup et al. [2012]. To use LDA for model choice with the programs discussed here is similar to the use of linear combinations for parameter inference in that the summary statistics of the observed and simulated data have to be transformed before running either of *ABCtoolbox* or the R package *abc* to perform model choice.

As an alternative to LDA, *ABCtoolbox* offers a greedy search algorithm to identify the combination of statistics having the largest power to discriminate between models. This search is done iteratively, by firstly evaluating the power of each single statistic and then adding additional statistics until no increase in power is observed. To perform this type of analysis, *ABCtoolbox* has to be run in the `findStatsModelChoice` mode and by providing the simulation files for at least two models, as well as the parameters required to perform the estimation. In addition, and using the argument `modelChoiceValidation`, one also needs to specify the number of simulations to be used as pseudo-observed data in each iteration to evaluate the power. As an example, consider the following input file for performing this type of analysis to find summary statistics appropriate for contrasting the normal versus uniform distribution toy models.

```
task findStatsModelChoice
simName simNorm.txt ;simUnif.txt
obsName simple.obs
maxCor 1.0
params 1-2;1-2
numRetained 1000
maxReadSims 100000
outputPrefix ABC_searchStats
modelChoiceValidation 1000
```

The results of this analysis are written to a file with tag `searchStatsgreedySearch`. Part of this file is shown in Table 7. As shown there, the power to distinguish between these models is very high for multiple sets of summary statistics. Generally, it is recommended to choose the smallest among all sets with highest power, which would be the set consisting of the statistics `var` and `range`. As is shown in Figure 2, the two-dimensional distribution of these two statistics is indeed rather different between the models.

4.4 Generating simulations

For simple models such as the normal distribution example that we examined in the previous section, it is relatively easy to perform the simulations using custom scripts written in scripting languages such as R. However, for realistically complex models we often rely on specialized programs for performing simulations. Moreover, for certain ABC variants such as ABC-MCMC, the simulation and estimation procedures are inherently linked, thus requiring running the program that performs simulations jointly with the program that performs ABC. In this section we will illustrate how to use the program *ABCtoolbox* as well as the R package *EasyABC* to automate and streamline the simulation process and to perform more sophisticated ABC techniques such as ABC-MCMC.

4.4.1 Generating simulations for rejection

We will first focus on how to use these two ABC pipelines to generate simulations from parameters values drawn from prior distributions. The such generated simulations are then ready to be used with all the estimation techniques introduced above.

EasyABC This R package allows the user to launch simulations from an external program and to retrieve the output of these simulations in a format ready for post-processing with the R package *abc* using the function `ABC_rejection()`. To achieve this, the user has to provide both a list containing the definitions of the prior distributions, as well as a model definition. The list containing prior distributions simply contains the names of the desired distributions, along with their arguments. For instance, a list defined as

```
prior <- list(c("unif", 500, 1000), c("norm", 10, 100));
```

will imply that there are two model parameters with the prior on the first being a uniform distribution bounded at 500 and 1000, and the prior on the second being a normal distribution with mean 10 and standard deviation 100.

The model may be either an R function taking the parameters as arguments and returning a vector of summary statistics, or the name of an executable (such as a BASH script) that will be used to generate the simulation. In case an executable is given, it is assumed that this executable will read the model parameters to be used from a file called `input` written by the R function, and write the resulting summary statistics to a file called `output` that, in turn, will be read by the R function. Since most programs do not follow this structure, it is often required to wrap them in a BASH script to interact with the *EasyABC* package properly. As an example, consider a script `executable.sh` that wraps a program to run the simulation of a model with two parameters. One can then use *EasyABC* to generate 10^4 simulations under this setting and using the prior definition shown above as follows:

```
ABC_sim <- ABC_rejection(model=binary_model('executable.sh'), prior=prior, nb_simul  
=10^4);
```

Table 8: Declaration of parameters in the *estName* file.

Column	Content
1	Indicator 1/0 for being integer or rational number.
2	Name of the parameter.
3	Type of prior (see <i>ABCtoolbox</i> Manual for the types of supported priors).
4 -	Parameters for prior (for example min,max for uniform prior).
Last	Indicator output/hide for whether to print the parameter in the output file.

ABCtoolbox An even more advanced and feature-rich way of using existing programs to generate simulations is offered by *ABCtoolbox*. To do so, *ABCtoolbox* has to be run in **simulate** mode, specified with the argument **task**. similarly to *EasyABC*, the user then needs to specify both the model parameters and their prior distributions, as well as how to use existing programs to generate simulations using values drawn from the prior.

The model parameters and their priors have to be provided through an external file referred to as the **est** file, the name of which is provided with the argument **estName**. This file is structured in three distinct sections called [PARAMETERS], [RULES] and [COMPLEX PARAMETERS]. Only the first of those is mandatory and contains the definitions of the model parameters for which estimations are to be carried out. These model parameters and their prior distributions are declared using multiple columns as explained in Table 8. In brief, the first column indicates whether or not a model parameter is to be truncated to an integer value, the second column lists the name of the parameter and the third column the prior distribution function. The remaining columns contain the parameters for this distribution, for instance the lower and upper bound, as wells as the mean and standard deviation for a normal prior. The last column specifies whether or not the parameter values are to be printed to the output file.

As an example, consider the following **est** file.

```
[PARAMETERS]
0 PARAM_A unif -1 1 output
0 PARAM_B norm -10 10 1 2 output
[RULES]
PARAM_A > PARAM_B
[COMPLEX PARAMETERS]
0 PARAM_B_SCALED = exp(PARAM_B) / PARAM_A
```

Here, we made use of the optional [RULES] section to limit the simulations to cases where **PARAM_A** is larger than **PARAM_B**. In addition, we benefited from [COMPLEX PARAMETERS] section to define a new variable **PARAM_B_SCALED**, which will always be set to the exponential of **PARAM_B**, divided by the value of **PARAM_A**. *ABCtoolbox* will understand most mathematical symbols and offers a wide variety of functions in this section, which allows for the definition of prior distributions and model parameterization in a different way than required by the simulation software.

To specify how *ABCtoolbox* is to interact with the external simulation software, the arguments **simProgram** and **simArgs** are used, where the former defines the name of the executable to be used,

and the latter the arguments to be passed to that executable. These arguments may contain tags referring to model parameters listed in the `est` file, as well as any other string. In case the simulation program reads the parameters from a specific input file, *ABCtoolbox* can be set up to scan such a file and to replace all occurrences of model parameter tags defined in the `est` file with their current values, and to save the result to a new file, which is then passed to the simulation program. To make use of this feature, the name of the input file has to be specified with the argument `simInputName`, and the tag `SIMINPUTNAME` may then be used to refer to the newly created input file among the arguments passed to the simulation program.

To perform simulations using the executable `myModel` taking two model parameters as arguments on the command line and using the `est` file defined above, an appropriate input file may look as follows:

```
task simulate
obsName sumstats.obs
estName Rules.est
numSims 100000
simProgram myModel
simArgs -a PARAM_A -b PARAM_B_SCALED -n 1
```

In case the simulation program is generating data instead of directly summary statistics itself, *ABCtoolbox* can run additional programs to do extra operations on the output of the simulation program, such as the calculation of summary statistics. Such a program can be defined with the argument `sumStatProgram` and the command-line arguments for the program are set with `sumStatArgs`. Note that `sumStatProgram` will always run after `simProgram`. A list of commonly used arguments when running *ABCtoolbox* in `simulate` mode are listed in Table 9.

4.5 Performing MCMC

Several other likelihood-free algorithms have been proposed that overcome the inherently low acceptance rates of rejection algorithms, among them a Markov chain Monte Carlo sampler (ABC-MCMC) first introduced by Marjoram et al. [2003] and sequential Monte Carlo or particle samplers [Sisson et al., 2007, Beaumont et al., 2009]. While both the R package *EasyABC* as well as *ABCtoolbox* offer both types of algorithms, we will focus here on the use of ABC-MCMC with these tools.

The basic idea of ABC-MCMC is to replace the likelihood ratio in the Hastings ratio of a classic MCMC by an acceptance-rejection step using some tolerance ϵ . Such an ABC-MCMC chain is then generating samples directly from $\mathbb{P}(\|S - S_{obs}\| < \epsilon | \theta)$, where θ is the vector of model parameters, S and S_{obs} the simulated and observed vectors of summary statistics, respectively, and $\|\cdot\|$ some distance measure in the summary statistics space. Such an algorithm was shown to require much less simulations than standard ABC methods to obtain equally good posterior estimates [Marjoram et al., 2003]. However, it turned out to be relatively tricky to tune this algorithm to perform properly since the acceptance rate of such an algorithm is directly given by the absolute likelihood, rather than the relative likelihood in standard MCMC. A result of this is that ABC-MCMC chains may easily get stuck in regions of low likelihood, requiring a careful choice of both the tolerance ϵ as well as the

Table 9: *ABCtoolbox settings for simulation*

Setting type	Setting	Description
Basic	task	Possible options simulate and estimate.
	samplerType	Possible sampler types are standard, MCMC, SuffStatMCMC.
	numSims	No. of simulations to perform.
	outName	Prefix for output files.
	estName	Filename for definitions of priors for parameters and rules.
	simProgram	Program to perform simulations.
	simArgs	Arguments for simulation program.
	obsName	File containing observed summary statistics.
	sumStatProgram	Program to be run after simProgram. For example a script calculating summary statistics.
	sumStatArgs	Arguments for sumStatProgram.
	sumStatName	File containing simulated summary statistics.
	doBoxCox	Do boxcox transformation.
	linearCombName	File containing linear combinations for transformation of statistics. (e.g., PLS components).
	doBoosting	Use all product combinations of statistics as additional statistics.
MCMC	numCaliSims	No. of calibration simulations.
	thresholdProp	Tolerance proportion of calibration simulations.
	rangeProp	Range of proposals.
	startingPoint	Starting location set from a random simulation (random) or the simulation with the minimum distance to the observed data (best).
	mcmcSampling	Interval between iterations that are printed in the output file.

initial starting positions. To improve the performance of this algorithm, we have proposed to tune the ABC-MCMC algorithm by means of an initial training set of simulations [Wegmann et al., 2009], which has been adopted by both *EasyABC* as well as *ABCtoolbox*. Specifically, the idea of such a calibration step is to choose a tolerance value ϵ that will result in sufficiently high acceptance rates and to find starting values in high likelihood regions. As with the classic rejection algorithm, it may be useful to transform summary statistics when calculating distances [Wegmann et al., 2009], and hence both *EasyABC* as well as *ABCtoolbox* offer to specify such transformations to be used during an ABC-MCMC chain.

While generally faster, an important issue with ABC-MCMC as well as Sequential Monte Carlo algorithms is that their output can not be directly used for validation. Instead, validation has to be done by repeating the whole process using pseudo-observed data sets, which may easily eat away the computational benefit of using these methods.

EasyABC In *EasyABC*, the ABC-MCMC algorithm is offered through the function `ABC_mcmc()`, which takes similar arguments as the function to perform the rejection algorithm, namely a list with prior definitions as well as a model, but also requires the vector containing the observed summary statistics to be specified using the argument `summary_stat_target`. In addition, several arguments for tuning the actual MCMC run are required. As an example, consider the following R code to generate posterior samples using ABC-MCMC for our normal toy model, using the function `calc.stats()` and the vector of observed summary statistics `S.obs` introduced above.

```
#define model
toy_model <- function(x){
  data <- rnorm(100, x[1], sqrt(x[2]));
  return(calc.stats(data));
}
toy_prior <- list(c("unif", -1, 1), c("unif", 0.1, 3));

#run ABC-MCMC
ABC_posterior <- ABC_mcmc(method="Wegmann", model=toy_model, prior=toy_prior, n_
  between_sampling=1, n_rec=10000, summary_stat_target=S.obs, n_calibration=10000,
  tolerance_quantile=0.01, numcomp=2);
```

Here, the argument `n_rec` specifies that 10,000 samples are to be generated. Further, the arguments `n_calibration` and `tolerance_quantile` specify that the ABC-MCMC chain will be calibrated from 10,000 simulations conducted under the prior, of which a fraction of 0.01 will be retained to calibrate the MCMC chain. Finally, the argument `numcomp` specifies that the total set of summary statistics is to be transformed into 2 PLS components.

Since an ABC-MCMC run is generating posterior samples, the output can be used directly to plot posterior distributions.

```
par(pty="s", mfrow=c(1, 2));
plot(density(ABC_posterior$param[, 1], from=-1, to=1, adjust=3), main="", xlab=expression(
  mu));
```

Table 10: *Downsampled synonymous SFS.*

Site class	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24
Site count	9906	7	5	2	0	1	1	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	77

```
plot(density(ABC_posterior$param[,2], from=0.1, to=3, adjust=3), main="", xlab=expression(
  sigma^2));
```

ABCtoolbox To perform the ABC-MCMC algorithm with *ABCtoolbox*, a few arguments have to be added to the input file shown above for standard sampling. First, the argument `samplerType` has to be set to `MCMC`. Then, the arguments `numCaliSims`, `thresholdProp` and `rangeProp` are used to specify the number of simulations to be used for calibration, the fraction of those simulations to be used to calibrate the threshold, and the fraction of the standard deviation of parameter values among these retained simulations to be used to propose new values during the MCMC chain, respectively. To transform the summary statistics during the MCMC chain, a file with the definition of linear combinations can be provided with the argument `linearCombName`. To use PLS transformations, for instance, an initial set of calibration simulations can be used to find appropriate PLS components as discussed above, and the resulting PLS definition file is then provided using this argument. For an example of an input file we refer the reader to the population genetics example discussed below.

4.6 A population genetics example

Here we will illustrate how to implement techniques described in the previous sections to estimate important aspects of the recent human demographic history from an allele frequency data set made publicly available by Boyko et al. [2008]. Specifically, we will use the site-frequency spectrum (SFS) for synonymous sites obtained for a sample of 24 African Americans (from Table S2 in Boyko et al. [2008]) to infer the parameters of a simple population genetic model in which the ancestral African population was of size N_{ANC} , but experienced an instantaneous change in size t generations ago to N_{CUR} . Following Boyko et al. [2008], we parameterized the time of the size change in units of the current population size ($\tau = t/(2 \times N_{CUR})$) and to allow the simulations to be performed in a time reasonable for an illustrative example, we downsampled the original data from the original 5 million sites to the SFS of only 10,000 sites shown in Table 10. From this data, we then calculated the set of classic population genetic summary statistics shown in Table 11. These summary statistics are then stored in the file `popgen.obs` for later usage.

The goal is then to use *ABCtoolbox* to generate posterior samples under this model using an ABC-MCMC chain, and we will provide a detailed step-by-step guidance on how to achieve that. The first step always consists of defining the model parameters in the `est` file. For the model concerned here, we will use the `est` file `popgen.est` provided below.

Rules file

```
[PARAMETERS]
0 LOG10_N_CUR    unif 2 6    output
```

Table 11: Summary statistics for synonymous sites.

Statistic	Header tag	Value
No. of singletons	sfs1	7
No. of segregating sites (S)	S	17
Average pairwise diversity (π)	pi	3.061594
Waterson's thita	thita	4.552403
Tajima's D	taj_D	-1.174098

```

0 LOG10_OMEGA    unif    -3    3    output
0 TAU            unif 0 1    output
[COMPLEX PARAMETERS]
1 N_CUR = pow10(LOG_N_CUR)    hide
1 T1 = TAU * 2 * NPOP    hide
0 OMEGA = pow10(LOG10_OMEGA)    hide

```

As can be seen from this file, we decided to put uniform priors on the logarithm of the current population size N_{CUR} and the relative size of the ancestral population $\omega = \frac{N_{ANC}}{N_{CUR}}$, but a uniform prior on the relative age of the population size change τ .

To generate simulations under this model, we will make use of the program *fastsimcoal2* that allows to simulate SFS under various demographic scenarios [Excoffier et al., 2013]. However, *fastsimcoal2* requires the parameters to be specified differently from our priors, and we thus make use of the [COMPLEX PARAMETERS] section to transform our model parameters appropriately. Specifically, we have to provide N_{CUR} and the population size change on the natural scale, and further the age of the population size changes in generations. We then prepare the input file `popgen.par` for *fastsimcoal2* that specifies this model, using the parameter tags defined in the input file. While explaining the details of how to use *fastsimcoal2* for such a model is beyond the scope of this chapter, we provide the corresponding input file in the appendix and refer the reader to the manual of *fastsimcoal2* for more details. To calculate summary statistics from the simulated data we will use the custom perl script `calcPopstats.pl` also provided in the appendix to this chapter.

In order to use PLS transformation during the MCMC chain, we first generated an initial set of 1000 simulations using *ABCtoolbox* using the following input file.

```

task simulate
obsName popgen.obs
estName popgen.est
numSims 1000
outName popgen
simInputName popgen.par
simProgram fastsimcoal2
simArgs -i 1popgen-temp.par -s -d -n 1 -q -x
sumStatProgram calcPopstats.pl
sumStatArgs popgen-temp/popgen-temp_DAFpop0.obs
doBoosting

```

Here we specify how *ABCtoolbox* is interacting with *fastsimcoal2* to generate simulations as well as with our custom perl script to calculate summary statistics from the output. While we refer the reader to the manual of *fastsimcoal2* for the details in the command line used, we note that the output written by *fastsimcoal2* will be located in a subdirectory and have a specific name. We thus provide the path to this file to our perl script `calcPopstats.pl` via command line arguments. In contrast to previously discussed input files we also added the additional argument `doBoosting`, which will tell *ABCtoolbox* to also add all squares and pair-wise products of calculated summary statistics as additional summary statistics. This often proves helpful in exploiting non-linear relationships between parameters and statistics when finding linear combinations.

PLS components are then readily identified by following the steps discussed in section 4.3.1. By looking at the RMSE plot we found that 4 PLS components are sufficient to capture the information contained in the total of 20 summary statistics (included the boosted ones). Having the appropriate PLS definition file `PLSdef_popgen.txt` at hand, we can then set up *ABCtoolbox* to run an ABC-MCMC chain using the following input file.

```
task simulate
samplerType MCMC
obsName popgen.obs
estName popgen.est
numSims 1000
outName popgen
simInputName popgen.par
simProgram fastsimcoal2
simArgs -i 1popgen-temp.par -s -d -n 1 -q -x
sumStatProgram calcPopstats.pl
sumStatArgs popgen-temp/popgen-temp_DAFpop0.obs
doBoosting
numCaliSims 1000
thresholdProp 0.1
rangeProp 1
linearCombName PLSdef_popgen.txt
numLinearComb 4
doBoxCox
```

This input file differs from the previous one in that the argument `samplerType` was added to tell *ABCtoolbox* to run an ABC-MCMC chain, and in that the arguments required for the calibration step (`numCaliSims`, `thresholdProp` and `rangeProp`), and those to use linear combinations of summary statistics (`linearCombName`, `numLinearComb` and `doBoxCox`) were added. Note that the script to find linear combinations provided by *ABCtoolbox* performs a Box-Cox transformation on the summary statistics, and hence in order to use the generated PLS definition file, *ABCtoolbox* needs to perform a similar transformation first, which is requested with the argument `doBoxCox`.

While the output of such a run already corresponds to samples taken from the posterior distribution $\mathbb{P}(\|S - S_{obs}\| < \epsilon | N_{CUR}, \omega, \tau)$, an additional improvement may be achieved by conducting an ABC-GLM estimation with an additional rejection step that will further reduce the threshold ϵ and hence




Figure 7: Marginal posterior estimate for parameter $\log_{10}N_{CUR}$ (left panel) and joint posterior estimate for parameters τ and $\log_{10}(N_{ANC}/N_{CUR})$ (right panel). Solid contour lines specify high posterior density intervals.

the accuracy of the posterior. Such an analysis is conducted as described in Section 3.2 and the resulting posteriors may then be plotted in R.

The posterior estimates we obtained for parameters N_{CUR} , ω and τ (Figure 7) indicated a large expansion that happened ~ 13000 generations ago to a present effective population size of ~ 32000 . The credible intervals of the posterior estimates are large due to the small size of the downsampled data set. However, the estimates are in good agreement with the findings of the Boyko et al. [2008], who used a maximum likelihood method on the full data.

References

- Simon Aeschbacher, Mark a Beaumont, and Andreas Futschik. A Novel Approach for Choosing Summary Statistics in Approximate Bayesian Computation. *Genetics*, 192(November):1027–1047, September 2012a. ISSN 1943-2631. doi: 10.1534/genetics.112.143164.
- Simon Aeschbacher, Mark A. Beaumont, and Andreas Futschik. A novel approach for choosing summary statistics in approximate bayesian computation. *Genetics*, 192(3):1027–1047, November 2012b. ISSN 0016-6731, 1943-2631. doi: 10.1534/genetics.112.143164.
- Simon Barthelmé and Nicolas Chopin. ABC-EP: expectation propagation for likelihood-free Bayesian computation. In *Proceedings of the 28th international conference on machine learning*, Bellevue, WA., 2011.
- M. A. Beaumont, J.-M. Cornuet, J.-M. Marin, and C. P. Robert. Adaptive approximate Bayesian computation. *Biometrika*, 96(4):983–990, October 2009. ISSN 0006-3444. doi: 10.1093/biomet/asp052.
- Mark A. Beaumont. Joint determination of topology, divergence time, and immigration in population trees. In S. Matsumura, P. Forster, and C. Renfrew, editors, *Simulation, Genetics and Human Prehistory*. McDonald Institute for Archaeological Research, United Kingdom, 2008.
- Mark A. Beaumont, Wenyang Zhang, and David J. Balding. Approximate Bayesian Computation in Population Genetics. *Genetics*, 162:2025–2035, December 2002.
- M. G. B. Blum and O. François. Non-linear regression models for approximate bayesian computation. *Statistics and Computing*, 20(1):63–73, 2010. ISSN 0960-3174. doi: 10.1007/s11222-009-9116-0.
- M. G. B. Blum, M. A. Nunes, D. Prangle, and S. A. Sisson. A comparative review of dimension reduction methods in approximate bayesian computation. *Statist. Sci.*, 28(2):189–208, 05 2013. doi: 10.1214/12-STS406.
- Adam R Boyko, Scott H Williamson, Amit R Indap, Jeremiah D Degenhardt, Ryan D Hernandez, Kirk E Lohmueller, Mark D Adams, Steffen Schmidt, John J Sninsky, Shamil R Sunyaev, Thomas J White, Rasmus Nielsen, Andrew G Clark, and Carlos D Bustamante. Assessing the evolutionary impact of amino acid mutations in the human genome. *PLoS Genetics*, 4(5):e1000083, May 2008. ISSN 1553-7404. doi: 10.1371/journal.pgen.1000083.
- T. C. Bray, V. C. Sousa, B. Parreira, M. W. Bruford, and L. Chikhi. 2BAD: an application to estimate the parental contributions during two independent admixture events. *Molecular Ecology Resources*, 10(3):538–541, 2010.
- Jui-Hua Chu, Daniel Wegmann, Chia-Fen Yeh, Rong-Chien Lin, Xiao-Jun Yang, Fu-Min Lei, Cheng-Te Yao, Fa-Sheng Zou, and Shou-Hsien Li. Inferring the geographic mode of speciation by contrasting autosomal and sex-linked genetic diversity. *Molecular biology and evolution*, 30(11):2519–30, November 2013. ISSN 1537-1719. doi: 10.1093/molbev/mst140.

- Jean-Marie Cornuet, Filipe Santos, Mark A. Beaumont, Christian P. Robert, Jean-Michel Marin, David J. Balding, Thomas Guillemaud, and Arnaud Estoup. Inferring population history with *DIY ABC*: a user-friendly approach to approximate Bayesian computation. *Bioinformatics*, 24(23): 2713–2719, 2008.
- Katalain Csilléry, Olivier François, and Michael G. B. Blum. abc: an R package for approximate Bayesian computation (ABC). *Methods in Ecology and Evolution*, 3:475–479, 2012.
- J. A. Cuesta-Albertos and A. Nieto-Reyes. The random tukey depth. *Computational Statistics & Data Analysis*, 52(11):4979–4988, July 2008. ISSN 0167-9473. doi: 10.1016/j.csda.2008.04.021.
- Arnaud Estoup, Eric Lombaert, Jean-Michel Marin, Thomas Guillemaud, Pierre Pudlo, Christian P Robert, and Jean-Marie Cornuet. Estimation of demo-genetic model probabilities with Approximate Bayesian Computation using linear discriminant analysis on summary statistics. *Molecular ecology resources*, 12(5):846–55, September 2012. ISSN 1755-0998. doi: 10.1111/j.1755-0998.2012.03153.x.
- Laurent Excoffier, Isabelle Dupanloup, Emilia Huerta-Sánchez, Vitor C. Sousa, and Matthieu Foll. Robust Demographic Inference from Genomic and SNP Data. *PLoS Genetics*, 9(10):e1003905, October 2013. ISSN 1553-7404. doi: 10.1371/journal.pgen.1003905.
- Nelson J. R. Fagundes, Nicolas Ray, Mark A. Beaumont, Samuel Neuenschwander, Francisco M. Salzano, Sandro L. Bonatto, and Laurent Excoffier. Statistical evaluation of alternative models of human evolution. *Proceedings of the Natural Academy of Sciences U.S.A.*, 104(45):17614–17619, November 2007.
- Paul Fearnhead and Dennis Prangle. Constructing summary statistics for approximate Bayesian computation: semi-automatic approximate Bayesian computation. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 74(3):419–474, June 2012. ISSN 13697412. doi: 10.1111/j.1467-9868.2011.01010.x.
- Matthieu Foll, Mark A. Beaumont, and Oscar Gaggiotti. An approximate Bayesian computation approach to overcome biases that arise when using amplified fragment length polymorphism markers to study population structure. *Genetics*, 179:927–939, 2008.
- Olivier François and Laval Guillaume. Deviance information criteria for model selection in approximate Bayesian computation. *Statistical Applications in Genetics and Molecular Biology*, 10(1):1–25, 2011.
- Franck Jabot, Thierry Faure, and Nicolas Dumoulin. EasyABC: performing efficient approximate bayesian computation sampling schemes using r. *Methods in Ecology and Evolution*, 4(7):684–687, 2013. ISSN 2041-210X. doi: 10.1111/2041-210X.12050.
- Christoph Leuenberger and Daniel Wegmann. Bayesian computation and model selection without likelihoods. *Genetics*, 184(1):243–252, January 2010. ISSN 0016-6731, 1943-2631. doi: 10.1534/genetics.109.109058.

- Juliane Liepe, Chris Barnes, Erika Cule, Kamil Erguler, Paul Kirk, Tina Toni, and Michael P. H. Stumpf. ABC-SysBio—approximate Bayesian computation in Python with GPU support. *Bioinformatics*, 26(14):1797–1799, 2010.
- Joao S. Lopes, David Balding, and Mark A. Beaumont. PopABC: a program to infer historical demographic parameters. *Bioinformatics*, 25(20):2747–2749, 2009.
- Paul Marjoram, John Molitor, Vincent Plagnol, and Simon Tavaré. Markov chain monte carlo without likelihoods. *Proceedings of the National Academy of Sciences*, 100(26):15324–15328, December 2003. ISSN 0027-8424, 1091-6490. doi: 10.1073/pnas.0306899100.
- Pavlos Pavlidis, Stefan Laurent, and Wolfgang Stephan. msABC: a modification of Hudson’s ms to facilitate multi-locus ABC analysis. *Molecular Ecology Resources*, 10(4):723–727, 2010.
- Benjamin M Peter, Daniel Wegmann, and Laurent Excoffier. Distinguishing between population bottleneck and population subdivision by a Bayesian model choice procedure. *Molecular ecology*, 19(21):4648–60, November 2010. ISSN 1365-294X. doi: 10.1111/j.1365-294X.2010.04783.x.
- Jonathan K. Pritchard, Mark T. Seielstad, Anna Perez-Lezaun, and Marcus W. Feldman. Population growth of human Y chromosome: a study of Y chromosome microsatellites. *Molecular Biology and Evolution*, 16(12):1791–1798, 1999.
- Christian P. Robert, Jean-Marie Cornuet, Jean-Michel Marin, and Natesh S. Pillai. Lack of confidence in approximate bayesian computation model choice. *Proceedings of the National Academy of Sciences*, 108(37):15112–15117, September 2011. ISSN 0027-8424, 1091-6490. doi: 10.1073/pnas.1102900108.
- SA A Sisson, Y Fan, and Mark MM Tanaka. Sequential Monte Carlo without likelihoods. *Proceedings of the National Academy of Sciences of the United States of America*, 104(6):1760–5, February 2007. ISSN 0027-8424. doi: 10.1073/pnas.0607208104.
- David A. Tallmon, Ally Koyuk, Gordon Luikart, and Mark A. Beaumont. ONeSAMP: a program to estimate effective population size using approximate Bayesian computation. *Molecular Ecology Resources*, 8:299–301, 2008.
- Simon Tavaré, David J. Balding, R. C. Griffiths, and Peter Donnelly. Inferring coalescence times from DNA sequence data. *Genetics*, 145:505–518, February 1997.
- M. Tenenhaus, J.P. Gauchi, and C. Ménardo. Régression PLS et applications. *Revue de statistique appliquée*, 43(1):7–64, 1995.
- Kevin R. Thornton. Automating approximate Bayesian computation by local linear regression. *BMC Genetics*, 10:35, 2009.

Daniel Wegmann, Christoph Leuenberger, and Laurent Excoffier. Efficient approximate bayesian computation coupled with markov chain monte carlo without likelihood. *Genetics*, 182(4):1207–1218, August 2009. ISSN 0016-6731, 1943-2631. doi: 10.1534/genetics.109.102509. 00118 PMID: 19506307.

Daniel Wegmann, Christoph Leuenberger, Samuel Neuenschwander, and Laurent Excoffier. ABCtoolbox: a versatile toolkit for approximate bayesian computations. *BMC Bioinformatics*, 11(1):116, March 2010. ISSN 1471-2105. doi: 10.1186/1471-2105-11-116. 00095 PMID: 20202215.

5 Appendix

Here we provide additional files required to replicate our population genetics example. First, we provide the input file `popgen.par` for *fastsimcoal2* specifying the population genetics model used. Note that we decided to simulate the SFS using 10 independent loci with 1000 sites each.

fastsimcoal input file

```
//Number of population samples (demes)
1
//Population effective sizes (number of genes)
N_CUR
//Sample sizes
24
//Growth rates negative growth implies population expansion
0
//Number of migration matrices : 0 implies no migration between demes
0
//historical event: time, source, sink, migrants, new size, new growth rate,
//migr.matrix
1 historical events
T1 0 0 1 OMEGA 0 0
//Number of independent loci [chromosome]
10 0
//Per chromosome: Number of linkage blocks
1
//per Block: data type, num loci, rec. rate and mut rate
DNA 1000 0.00000 MUTRATE 0.33
```

Further, we provide the custom perl script `calcPopstats.pl` used to calculate summary statistics from site frequency spectra simulated with the program *fastsimcoal2*.

Perl script to calculate statistics from SFS

```
#!/usr/bin/perl -w
use strict;
#read fastsimcoal output SFS file
my $sfsfile=$ARGV[0];
open(FILE,"<",$sfsfile) or die "can't open SFS file";
open (OUT, ">","summary_stats-temp.txt") or die "can't open sum-stats file";
my ($firstline,$header,$sfsline)=(<FILE>,<FILE>,<FILE>);
#split sfsline into sfs
my @SFS=split /\t/, $sfsline;
my @stats;
#calculate stats
my ($sum,$S,$a1,$a2,$taj_D)=(0,0,0,0,0);
my $n=@SFS-2;
my ($b1,$b2)=((($n+1)/(3*($n-1)),2*($n**2+$n+3)/(9*$n*($n-1)));
#No. Segregating. sites S
for (my $i=1;$i<$n;$i++){
    $sum=$sum+$i*($n-$i)*@SFS[$i];
    $S=$S+@SFS[$i];
    ($a1,$a2)=($a1+1/$i,$a2+1/$i**2);
}
#Thita and pi
my ($thita,$pi)=($S/$a1,2*$sum/($n*($n-1)));
#Tajima's D
my ($c1,$c2)=($b1-1/$a1,$b2-($n+2)/($a1*$n)+$a2/($a1**2));
my ($e1,$e2)=($c1/$a1,$c2/($a1**2+$a2));
if ($S>0) {$taj_D=($pi-$S/$a1)/sqrt($e1*$S+$e2*$S*($S-1));}
#print out stats
@stats=($SFS[1],$S,$pi,$thita,$taj_D);
print OUT join("\t","sfs1","totS","pi","thita","taj_D"),"\n",join("\t",@stats,"\n");
close(FILE);close(OUT);system("rm $sfsfile");
```